

# Introduction to Reinforcement Learning

- A type of machine learning where an **agent** learns by interacting with the environment.
- Uses a system of **rewards and punishments** to guide decision-making.
- Inspired by **behavioral psychology**.

## Key Elements of RL

1. **Agent** – Learner or decision-maker.
2. **Environment** – The world where the agent interacts.
3. **State (S)** – A representation of the environment.
4. **Action (A)** – Possible moves the agent can take.
5. **Reward (R)** – Feedback from the environment.
6. **Policy ( $\pi$ )** – Strategy to choose actions based on the state.
7. **Value Function (V)** – Expected long-term reward for being in a state.
8. **Q-Value (Q)** – Expected reward for taking a specific action in a given state.

# Types of RL

## 1. Active RL

- The agent actively explores and takes actions to learn optimal policies.
- Uses **trial and error** to maximize rewards.
- Examples: **Q-learning, Deep Q Networks (DQN)**.
- **Example Scenario:** A robot trying different paths in a maze to maximize rewards.

## 2. Passive RL

- The agent **follows a fixed policy** and evaluates it.
- Learns the value of states under the given policy without exploring alternatives.
- Example Algorithms: **Policy Evaluation, Monte Carlo Methods**.
- **Example Scenario:** A self-driving car following a preset route while estimating travel time.

### 3. Model-Free RL

- The agent learns directly from interactions without a model of the environment.
- Examples: **Q-learning, SARSA.**
- **Use Case:** Video game AI improving through trial and error.

### 4. Model-Based RL

- The agent builds a model of the environment and plans actions based on it.
- Examples: **Dynamic Programming, Policy Iteration.**
- **Use Case:** A robot simulating possible moves before taking action.

## 5. Online RL

- The agent updates its knowledge continuously as it interacts with the environment.
- Example: **Stock trading AI learning in real time.**

## 6. Offline RL

- The agent learns from a pre-collected dataset without real-time interactions.
- Example: **Training autonomous cars on past driving data.**

## Q-learning

- A **model-free, off-policy** reinforcement learning algorithm.
- Updates Q-values based on the highest possible future reward.
- Uses the Bellman equation:
- $Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$
- where:
  - $\alpha$  = learning rate
  - $\gamma$  = discount factor
  - $r$  = reward
  - $\max Q(s', a')$  = maximum future reward estimate
- **Example Use Case:** Game AI improving performance by learning optimal moves.

## **SARSA (State-Action-Reward-State-Action)**

- A **model-free, on-policy** RL algorithm.
- Learns action-value function based on the agent's own policy.
- Update rule:
- $Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma Q(s', a') - Q(s, a)]$
- Difference from Q-learning: Uses the next action **a'** from the same policy instead of the max Q-value.
- **Example Use Case:** Learning in a **safe driving AI**, where it avoids risky actions based on its current policy.

Feature	Active RL	Passive RL
Policy	Learns optimal policy	Follows a fixed policy
Exploration	Yes	No
Decision-Making	Learns from actions taken	Evaluates given actions
Example	Q-learning	Policy Evaluation

### Applications of RL

- Gaming:** AlphaGo, Dota 2 AI.
- Robotics:** Autonomous robots learning to walk.
- Finance:** Stock trading strategies.
- Healthcare:** Personalized treatment recommendations.
- Self-driving Cars:** Learning to navigate safely.